

Deferred Shading

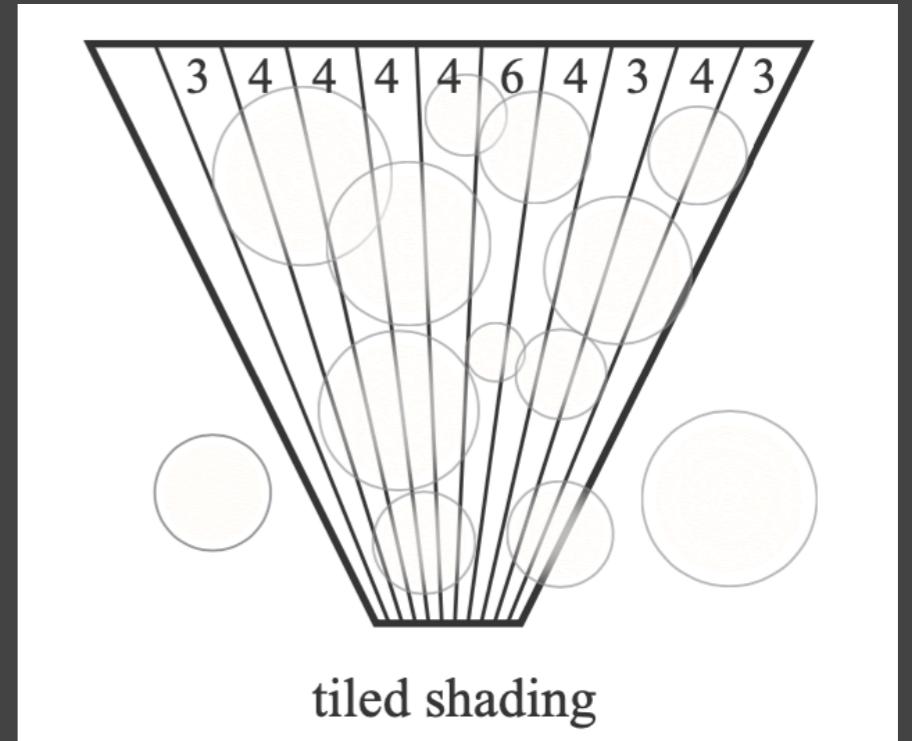
- Originally invented to **save shading time**
- Consider the rasterization process
 - Triangles -> fragments -> depth test -> shade -> pixel
 - Each fragment needs to be shaded (in what scenario?)
 - Complexity: $O(\#fragment * \#light)$
- Key observation
 - Most fragments will not be seen in the final image
 - Due to depth test / occlusion
 - Can we only shade those **visible fragments?**

Deferred Shading

- Modifying the rasterization process
 - Just **rasterize the scene twice**
 - Pass 1: no shading, just update the depth buffer
 - Pass 2 is the same (why does this guarantee shading visible frag. only?)
 - Implicitly, this is assuming **rasterizing the scene** is way faster than **shading all unseen fragments** (usually true)
 - Complexity: $O(\#fragment * \#light) \rightarrow O(\#vis. frag. * \#light)$
- Issue
 - Difficult to do anti-aliasing
 - But almost completely solved by TAA

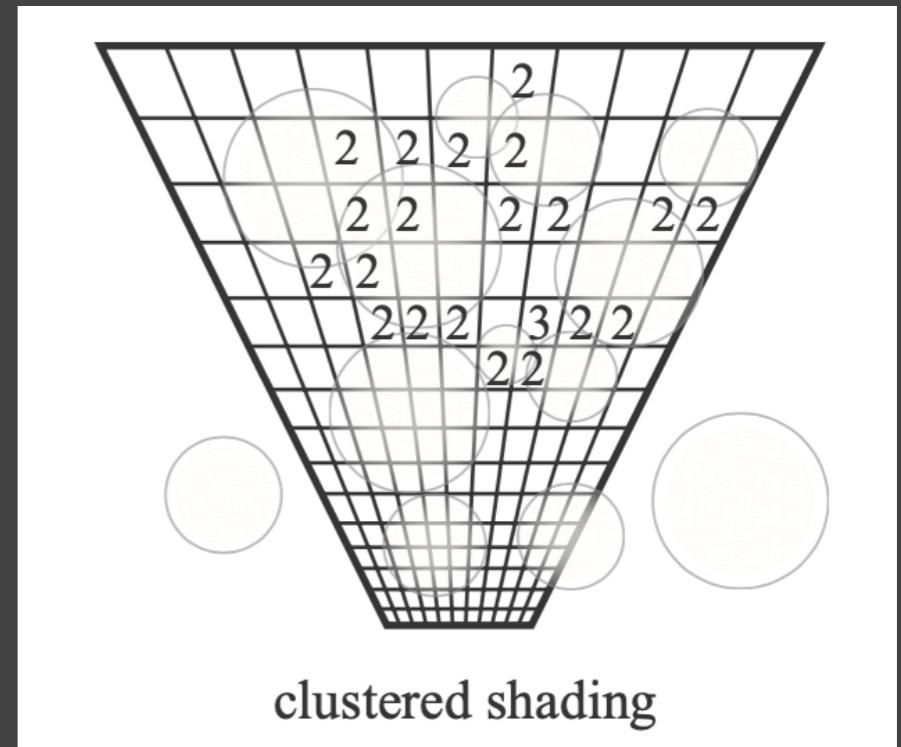
Tiled Shading

- Improvement: tiled shading
 - Subdivide the screen into tiles of e.g. 32x32 then shade each
- Key observation
 - Not all lights can illuminate a specific tile
 - Mostly due to the **square falloff with distance (!)**
 - Complexity: $O(\# \text{vis. frag.} * \# \text{light})$
-> $O(\# \text{vis. frag.} * \text{avg } \# \text{light per tile})$



Clustered Shading

- Further improvement: clustered shading
 - Further subdivide each tile into different depth segments
 - Essentially subdividing the view frustum into a 3D grid
- Key observation
 - The depth range of each tile can be quite large
 - Therefore, a lot of lights may be identified to have potential to lit the tile
 - But some lights may only lit a small depth range
 - Complexity: $O(\#vis. frag. * \text{avg } \#light \text{ per tile})$
-> $O(\#vis. frag. * \text{avg } \#light \text{ per cluster})$

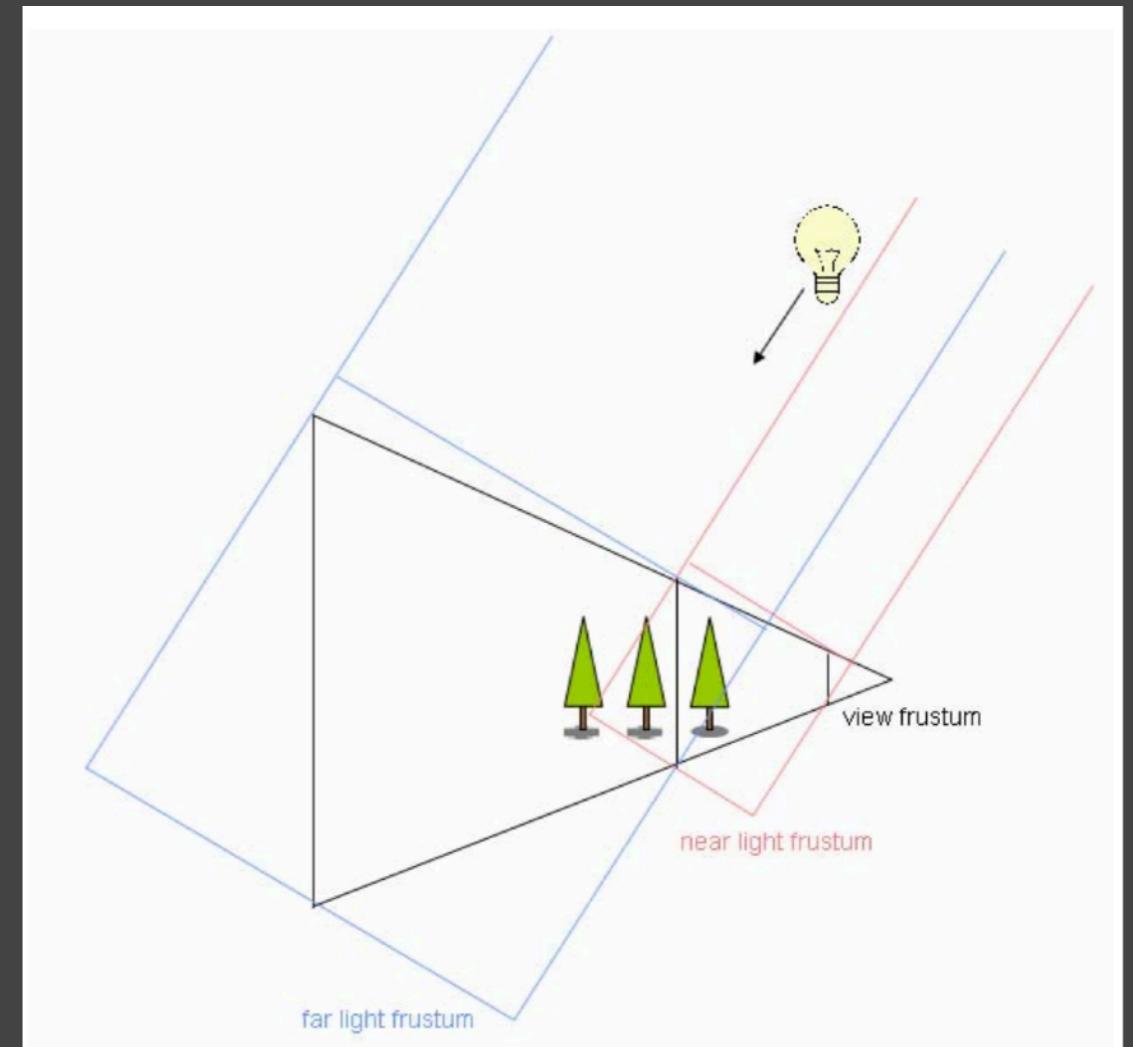
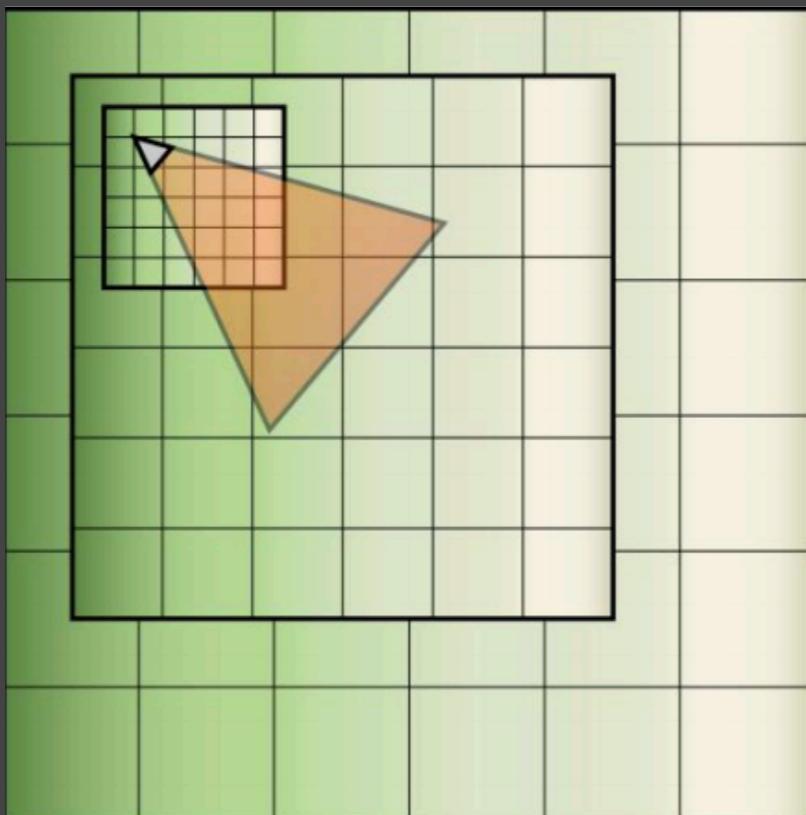


Level of Detail Solutions

- Level of Detail (LoD) is very important
 - Recall: texture MIPMAP-ing
 - Choosing the right level of detail to use can save computation
- The use of multiple levels of detail
 - Often called “cascaded” by the RTR industry

Level of Detail Solutions

- Example
 - Cascaded shadow maps
 - Cascaded LPV



[Dimitrov et al., Cascaded Shadow Maps]

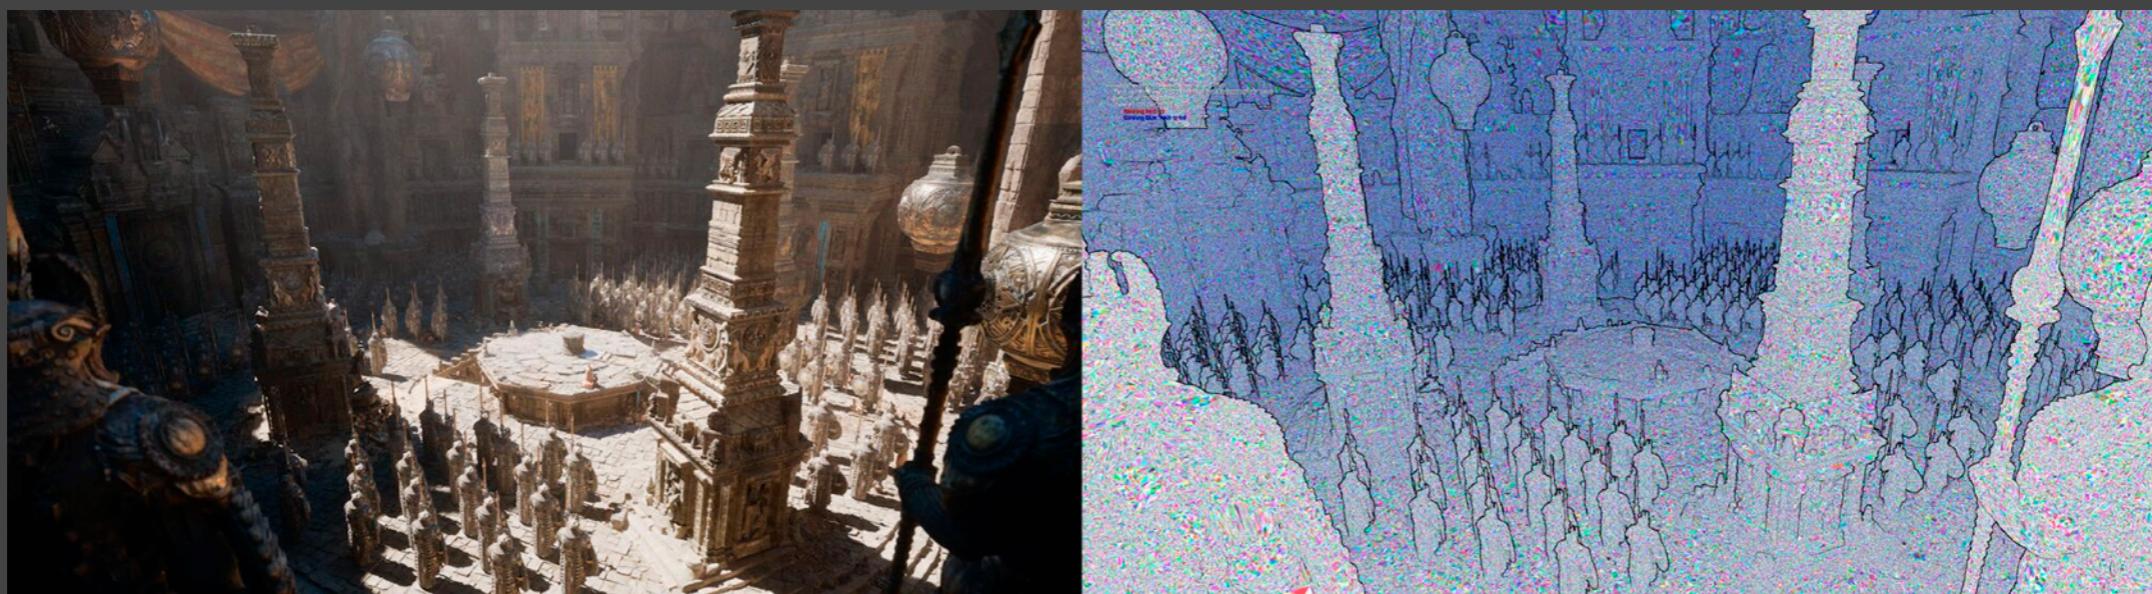
[Anton Kaplanyan, Light Propagation Volumes in CryEngine 3]

Level of Detail Solutions

- Key challenge
 - Transition between different levels
 - Usually need some overlapping and blending near boundaries
- Another example: geometric LoD
 - Recall: pre-generating a set of simplified obj. with different #tri.
 - Based on the distance to the camera, choose the right object to show (or part of obj., s.t. no triangle will be larger than a pixel)
 - Popping artifacts? Leave it to TAA!
 - This is Nanite in UE5 (but of course, Nanite has way more)

Level of Detail Solutions

- FYI, some (strongly) technical difficulties
 - Different places with different levels, how about cracks?
 - Dynamically load and schedule different levels, how to make the best use of cache and bandwidth, etc.?
 - Representing geometry using triangles or geometry textures?
 - Clipping and culling for faster performance?
 - ...



Global Illumination Solutions

- From this course, we can see that
 - Recall, when would screen space ray tracing (SSR) fail?
 - There is no single GI solution that is perfect for all cases, except for RRTT
 - But completely using RRTT is still too costly in the current generation
 - Therefore, the industry tends to use hybrid solutions
- For example, a possible solution to GI may include
 - SSR for a rough GI approximation (similar to our HW3)
 - Upon SSR failure, switching to more complex ray tracing
 - Either hardware (RRTT) or software (?)

Global Illumination Solutions

- Software ray tracing
 - HQ SDF for individual objects that are close-by
 - LQ SDF for the entire scene
 - RSM if there are strong directional / point lights
 - Probes that stores irradiance in a 3D grid
(Dynamic Diffuse GI, or DDGI)
- Hardware ray tracing
 - Doesn't have to use the original geometry, but low-poly proxies
 - Probes (RTXGI)
- The highlighted solutions are mixed to get Lumen in UE5

Summary: A Brief Q&A

- What is interesting?
 - Anything that requires thinking
 - Therefore, giving up thinking == committing suicide
- Is implementation less important than theory?
 - NEVER. But engineering skills must be acquired in engineering.
- You don't teach implementation, does it mean that you are not good at programming?
 - Dude, I was in Tsinghua's ACM/ICPC team

Questions?

Congratulations!



Real-time shadows / env. lighting



Real-time global illumination



Real-time shading / materials



Real-time ray tracing

Congratulations!

- Yet still, a lot of uncovered topics
 - Texturing an SDF
 - Transparent material and order-independent transparency
 - Particle rendering
 - Post processing (depth of field, motion blur, etc.)
 - Random seed and blue noise
 - Foveated rendering
 - Probe based global illumination
 - ReSTIR, Neural Radiance Caching, etc.
 - Many-light theory and light cuts
 - Participating media, SSSSS
 - Hair appearance
 - ...

Computer Graphics
is
AWESOME!

Advertisements

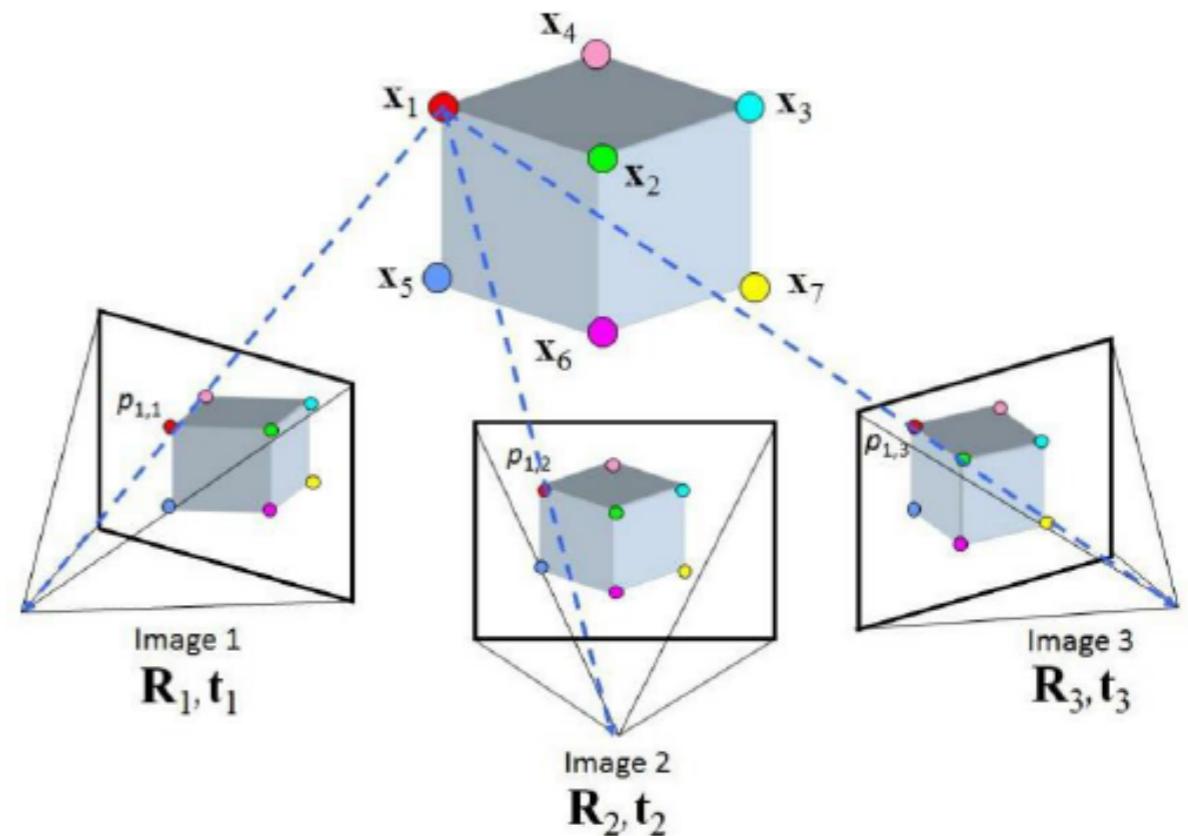
GAMES203

- 3D vision (reconstruction and more)
 - A combination of computer vision and computer graphics
- Starting July 2, 2021
- By Prof. Qixing Huang
 - From the University of Texas at Austin
- Let's take a glance at this course
 - 3 different topics



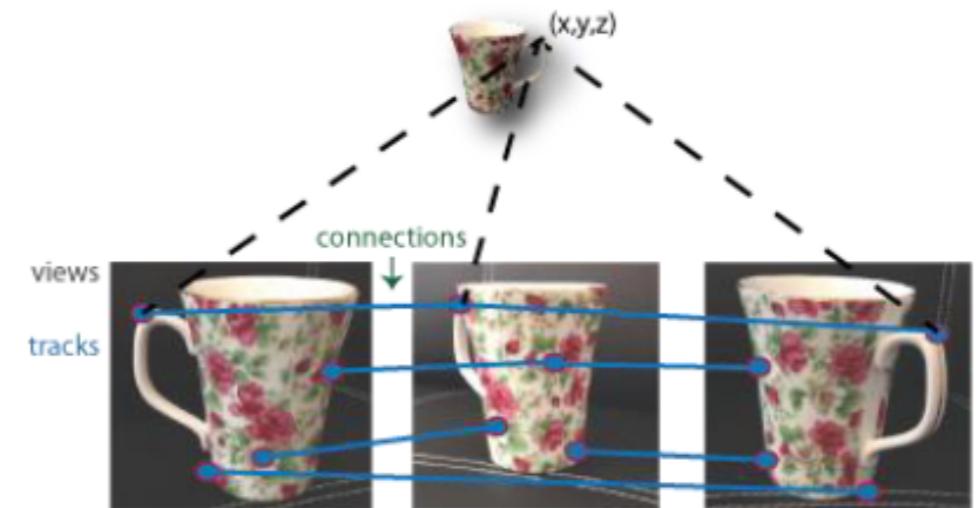
Topic I: 3D Reconstruction

- Geometry
 - Epipolar geometry
 - Fundamental matrix
 - Extrinsic/Intrinsic camera parameters
 - Camera calibration
 - Vanishing points
 - Homogeneous coordinates

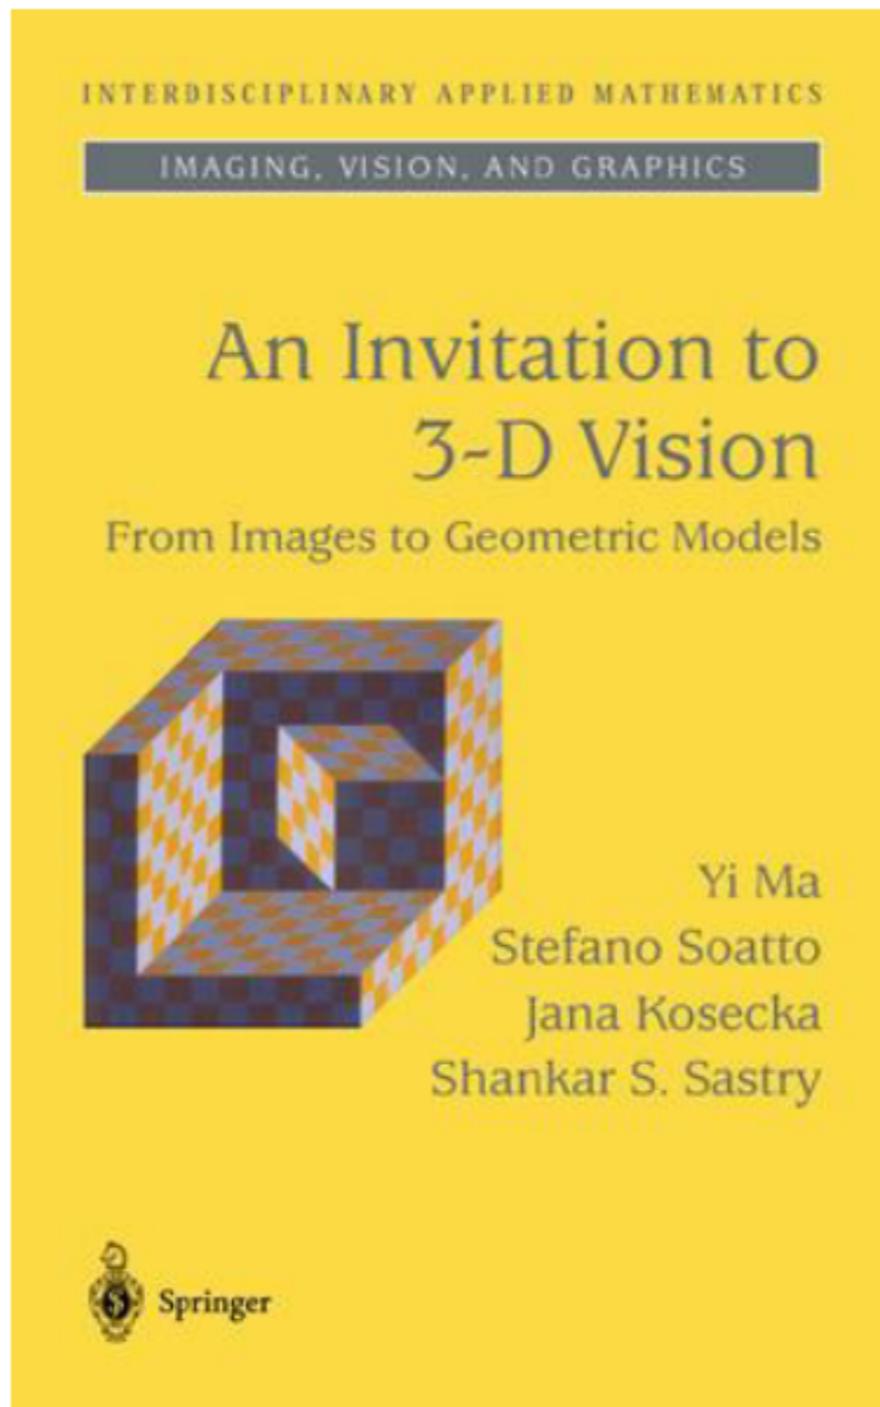


Topic I: 3D Reconstruction

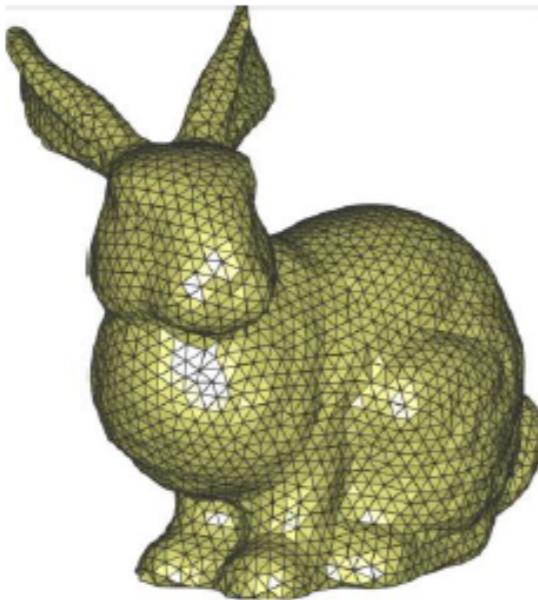
- Algorithms
 - Feature extraction
 - Feature correspondences
 - Relative camera pose
 - Structure-from-motion
 - Multiview stereo
 - Bundle adjustment
 - ICP



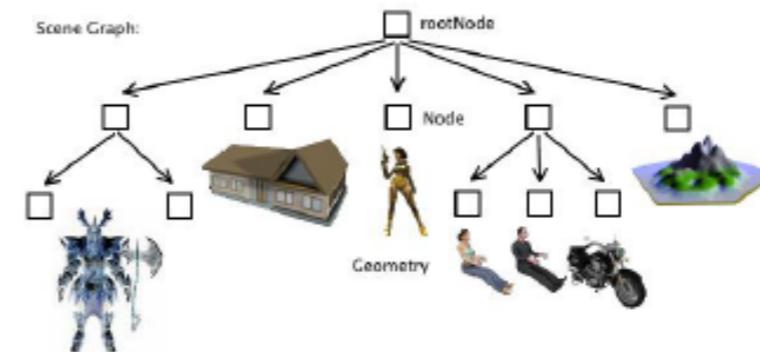
Textbook



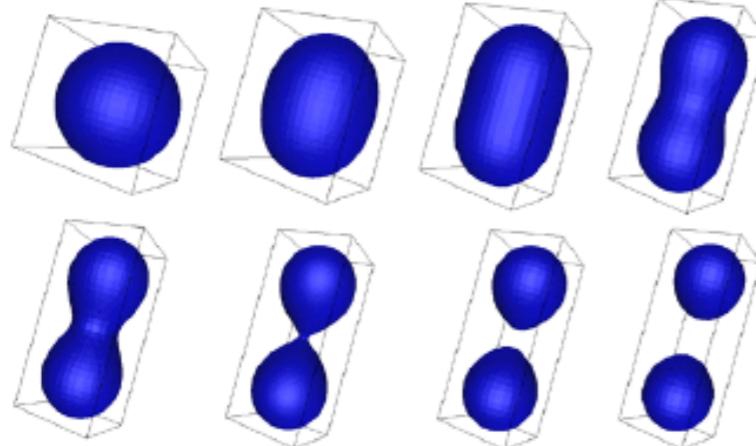
Topic II: How to represent 3D Data



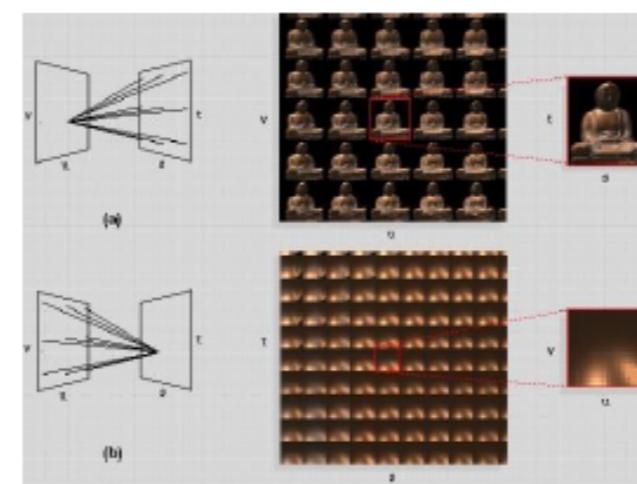
Triangular mesh



Part-based models



Implicit surface



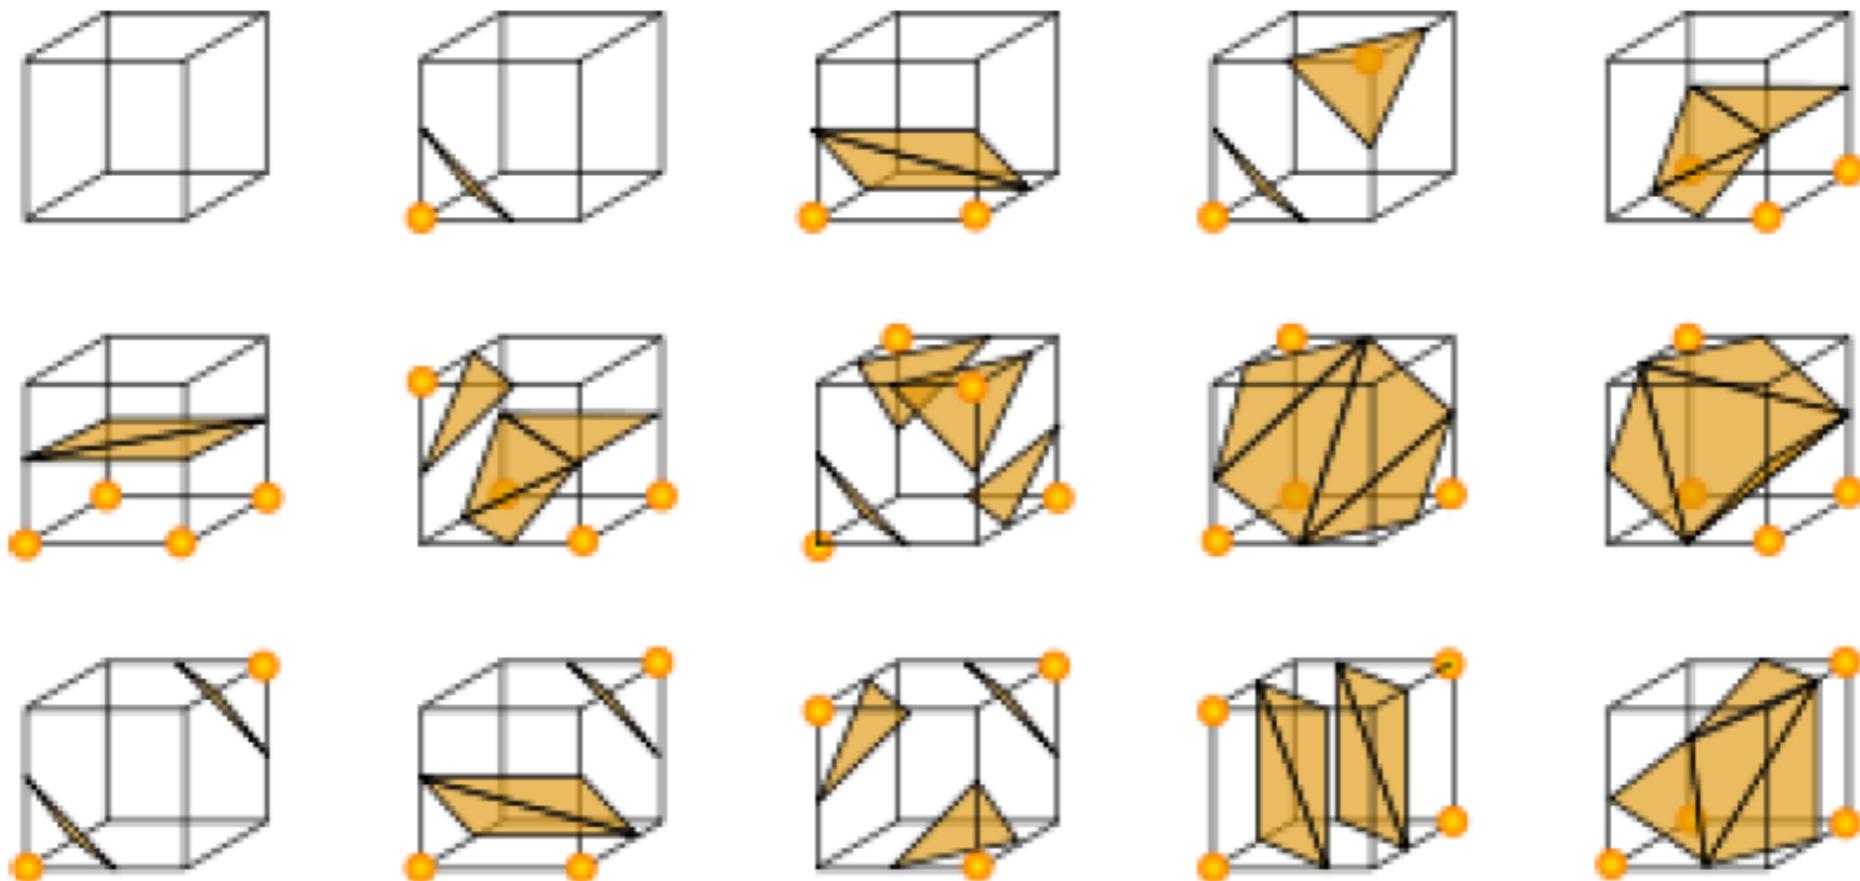
Light Field Representation



Point cloud

Conversion between different representations

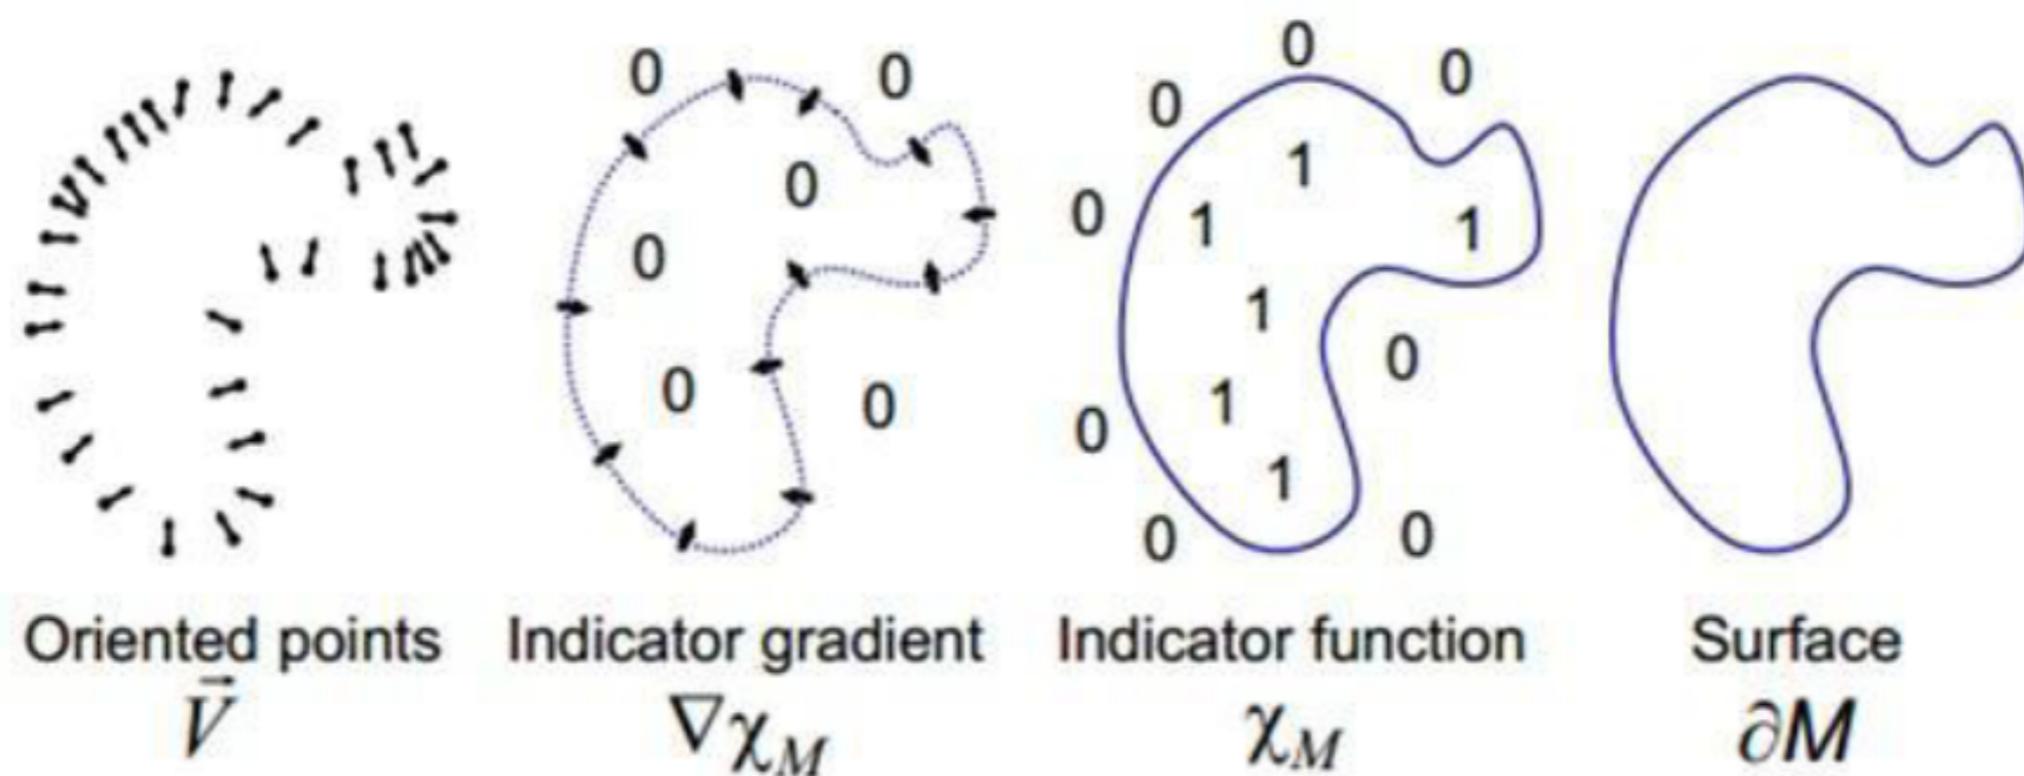
- Implicit -> mesh (Marching Cube)



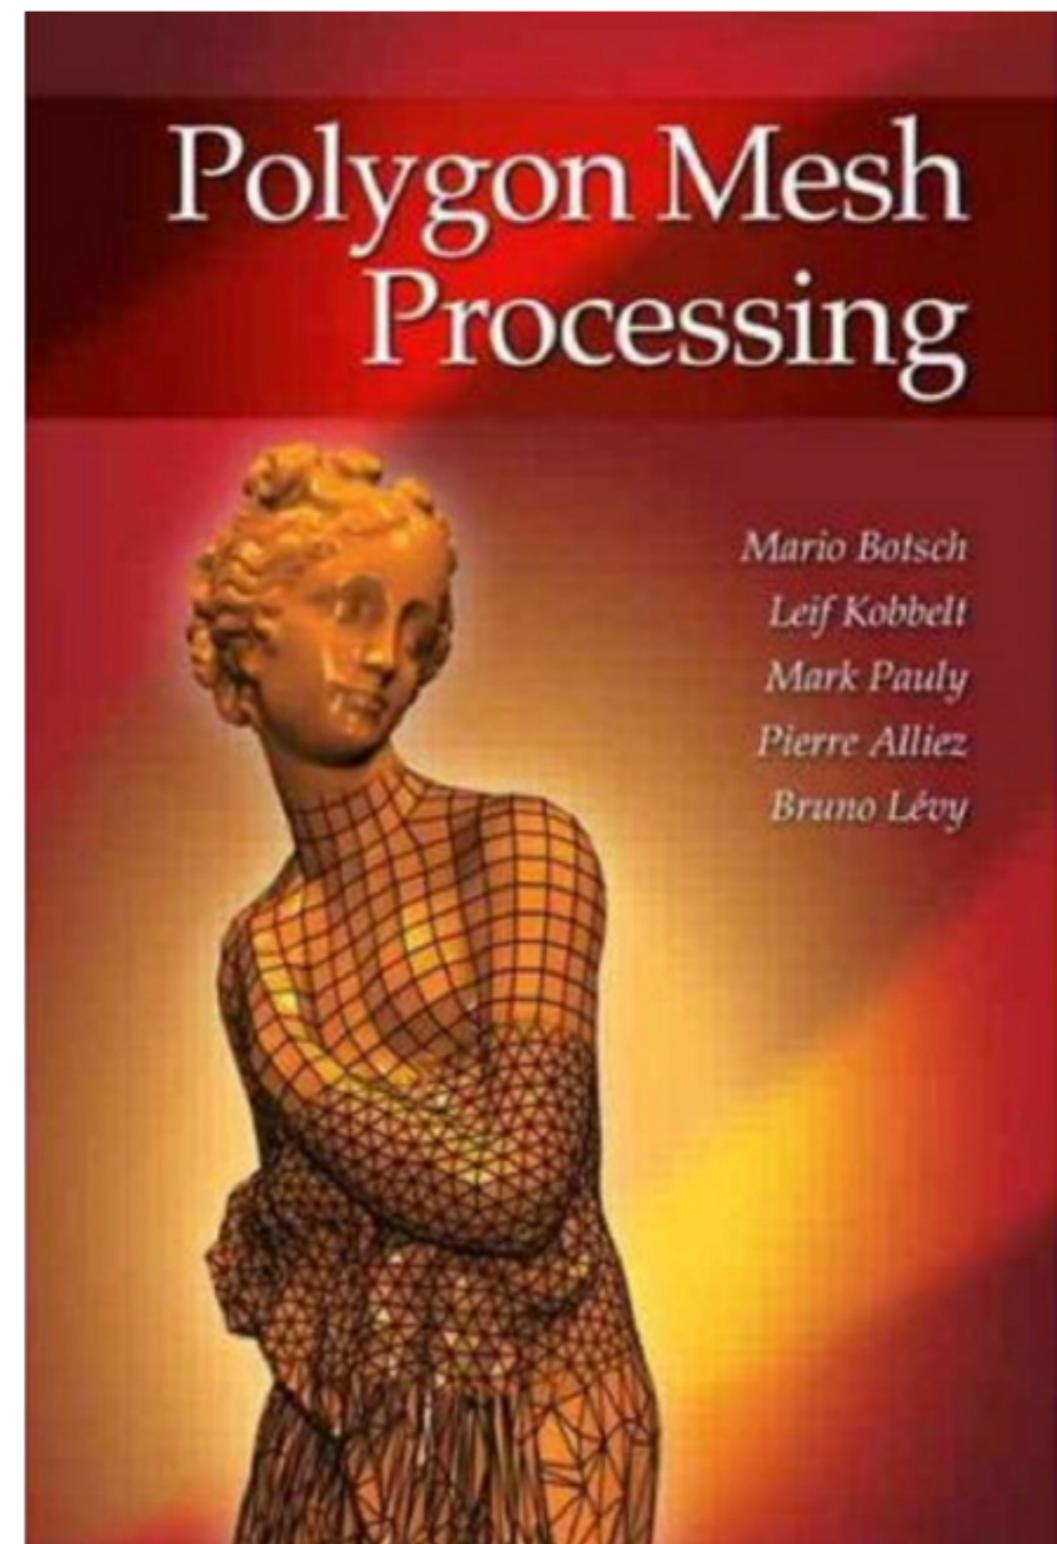
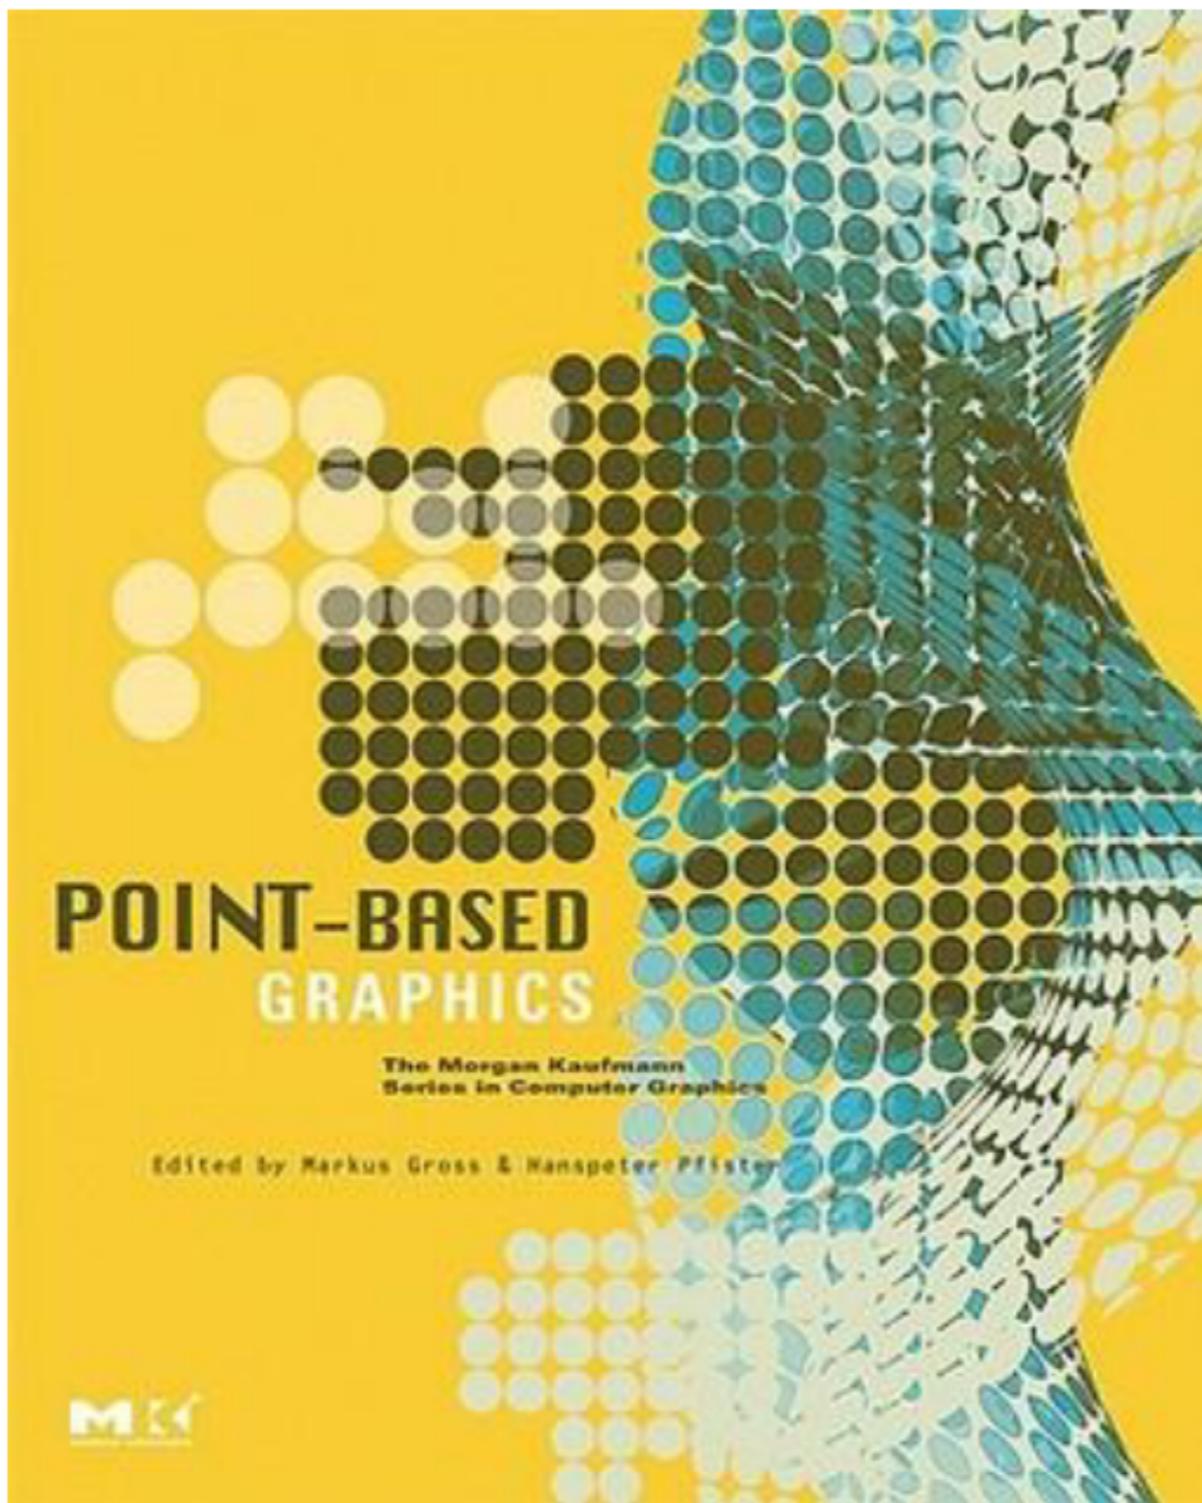
Conversion between different representations

- Pointcloud -> Implicit -> Mesh

[Kazhdan et al. 06]

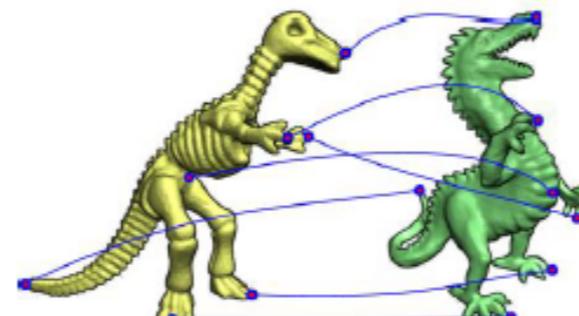
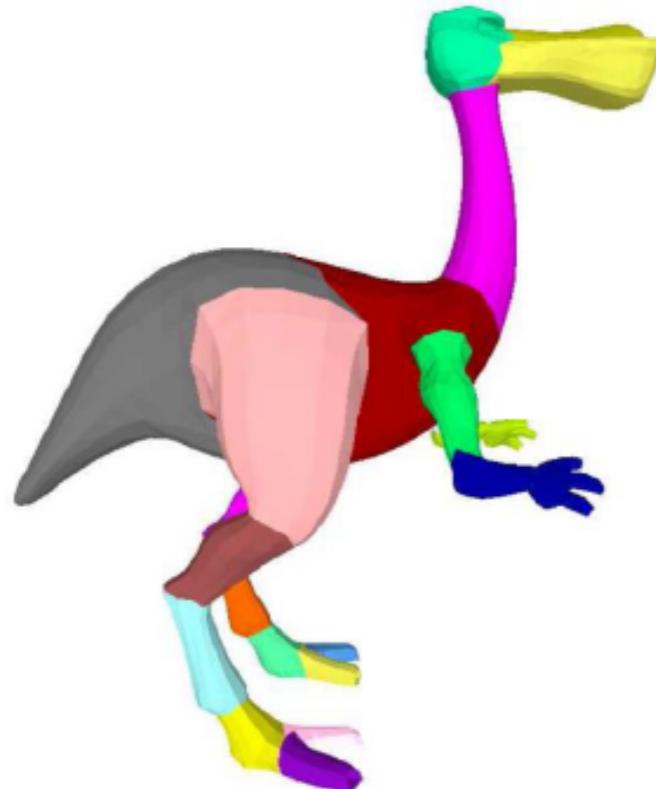


Two recommended books



Topic III: How to understand 3D Data

- Design algorithms to extract semantic information from one or a collection of shapes



Matching



[Funkhouser et al. 05]

Retrieval



[Mitra et al. 06]

Segmentation

Classification & Clustering

Another Rendering Course

- GAMES2XX
 - Unfortunately, GAMES3XX has been reserved for special topics
- Together with GAMES101 and GAMES202
 - A (hopefully helpful) computer graphics trilogy



Another Rendering Course

- GAMES2XX: Introduction to Offline Rendering / Advanced Image Synthesis
 - Part 1: Sampling and Light Transport
 - Part 2: Appearance Modeling
 - Part 3: State of the Art Research Topics
- Should be as **easy / comfortable / enjoyable** as this game 🐶 🐶 🐶

[Elden Ring, to appear]



GAMES: Graphics And Mixed Environment Symposium

图形学与混合现实在线平台

- 主页: <http://games-cn.org>
- 宗旨: 图形学及相关领域交流的华人[在线社区](#)
- 在线直播活动:
 - 每周四晚8:00-9:30的在线报告 (186期)
 - 专题: 几何、绘制、模拟、视觉、可视化...
 - 课程: 101 (闫令琪) 、 201 (胡渊鸣) 、 102 (刘利刚) 、 202 (闫令琪)
 - 已规划: 203 (黄其兴) 、 103 (王华明)
- 在线交流微信群: 16个群 (7900+人)



所有资料 (视频/PPT) 云端保存,
总观看 100+ 万人次

加入微信群的方法: 在微信中扫描右边的二维码, 加games
技术秘书为好友。然后回复“GAMES”即可获取群聊邀请。





Special Thanks to All of You!